Nominal AC-Matching

 $\begin{array}{c} \text{Mauricio Ayala-Rinc\'on}^{1[0000-0003-0089-3905]}, \text{ Maribel} \\ \text{Fern\'andez}^{2[0000-0002-1959-8730]}, \text{ Gabriel Ferreira Silva}^{1[0000-0003-1679-3597]}, \\ \text{Temur Kutsia}^{3[0000-0003-4084-7380]}, \text{ and Daniele} \\ \text{Nantes-Sobrinho}^{1,4[000-0002-1959-8730]} \end{array}$

University of Brasília, Brazil
 ayala@unb.br, and gabrielfsilva1995@gmail.com
 King's College London, U.K. maribel.fernandez@kcl.ac.uk
 Johannes Kepler University Linz, Austria kutsia@risc.jku.at
 Imperial College London, U.K dnantess@ic.ac.uk

Abstract. The nominal syntax is an extension of the first-order syntax that smoothly represents languages with variable bindings. Nominal matching is first-order matching modulo alpha-equivalence. This work extends a certified first-order AC-unification algorithm to solve nominal AC-matching problems. To our knowledge, this is the first mechanically-verified nominal AC-matching algorithm. Its soundness and completeness were verified using the proof assistant PVS. The formalisation enriches the first-order AC-unification algorithm providing structures and mechanisms to deal with the combinatorial aspects of nominal atoms, permutations and abstractions. Furthermore, by adding a parameter for "protected variables" that cannot be instantiated during the execution, it enables nominal matching. Such a general treatment of protected variables also gives rise to a verified nominal AC-equality checker as a byproduct.

Keywords: Nominal Matching \cdot Nominal AC-Matching \cdot Formal Methods \cdot PVS

1 Introduction

The nominal approach to the specification of systems with binders [20,23] extends first-order syntax with notions of name and binding that allow us to represent systems with binders smoothly. Such systems frequently appear in the formalisation of mathematics and when reasoning about the properties of programming languages. Taking into account α -equivalence is essential to represent bindings correctly. For example, the formulas $\forall x: x+1>0$ and $\forall y: y+1>0$ should be considered equivalent despite being syntactically different. Instead of using indices to represent bound variables, as in explicit substitution calculi à la de Bruijn, the nominal theory allows us to handle binders naturally, using atoms, atom permutations and freshness constraints to axiomatise α -equivalence [23,19].

Syntactic matching is the problem of, given terms t and s, finding a substitution σ such that $\sigma t = s$. Algorithms to solve matching problems are an essential component of functional languages and equational theorem provers: matching

is used to decide if an equation can be applied to a term. The problem of syntactic matching can be generalised to consider an equational theory E. In this case, called E-matching, we must find a substitution σ such that σt and s are equal modulo E, which we denote $\sigma t \approx_E s$. For example, if the system includes associative and commutative (AC) operators, such as + in the example above, then the matching algorithm should consider the AC axioms. Furthermore, equational programming languages, such as Maude, require efficient implementations of AC-matching to deal with AC-theories (see [16]).

If the system under study includes binders and AC operators, then α -equivalence should also be considered: for example, $\forall x: x+1>0$ should be considered equivalent to $\forall y: 1+y>0$. This paper focuses on the matching problem for languages that include binders and AC operators.

Nominal matching is the extension of first-order matching to the nominal syntax, replacing the notion of syntactic equality by α -equivalence. It has applications in rewriting, functional programming, and metaprogramming. For instance, various versions of matching modulo α -equivalence are used in functional programming languages that provide constructs for manipulating abstract syntax trees involving binders (e.g. [27,24]). In this work, we specify a nominal matching algorithm modulo AC function symbols (nominal AC-matching, for short) and prove its correctness and completeness using the proof assistant PVS.

Related Work. Nominal syntactic (i.e. modulo α -equivalence) equality-check, matching and unification were solved since the beginning of the development of the nominal approach; more than twenty years ago, Urban et al. [32] developed the first rule-based algorithm for nominal syntactic unification and further, Urban mechanised its correctness and completeness in Isabelle/HOL as part of the formalisation of the nominal approach in this proof assistant [30,31]. Furthermore, different approaches were designed to deal with nominal syntactic unification efficiently. Calvès and Fernández [12,11] and Levy and Villaret [21,22] developed polynomial nominal syntactic algorithms to solve unification problems and proved that nominal unification is a quadratic problem, respectively. Furthermore, Ayala-Rincón et al. [6] developed a nominal syntactic unification algorithm specified as a functional program and verified it in the proof assistant PVS. Enriching the nominal equational analysis with equational theories started with developing rule-based techniques for commutative operators. Such developments were initially checked in the proof assistant Coq and further in PVS [1,4]. Remarkable differences between nominal unification and nominal C-unification were discovered, such as the fact that when expressing solutions as pairs consisting of a freshness context and substitutions, nominal unification is unitary. In contrast, C-unification is not finitary [2,3].

Avoiding freshness constraints through a fixed-point approach was also studied as a mechanism to obtain finite complete sets of solutions [5]. Such fixed-point equations also appear in nominal techniques designed to deal with higher-order recursive let operators [25,26].

First-order AC-unification algorithms were proposed almost half a century ago, when Stickel [28,29] showed the connection between solving this problem

and computing solutions to linear Diophantine equations until a certain bound. Almost a decade later, Fages [17,18] fixed a mistake in Stickel's proof of termination. Since then, ideas to obtain more efficient AC-unification algorithms have been proposed, either by using a smaller bound when computing the solutions to the linear Diophantine equation [14], or by solving those equations more efficiently [14], or even by solving whole systems of linear Diophantine equations and using suitable data structures to represent the problem [10,8]. First-order AC-unification algorithms were not formalised until recently when a version of Fages' AC-unification algorithm was proved correct and complete using the proof assistant PVS [7]. This mechanisation applies the linear-Diophantine AC unification method discovered and fixed in works by Stickel and Fages [28,29,17,18], and can easily be adapted to deal with AC-equality and AC-matching problems as well. It is important to stress that such mechanisation was not a routine-formalisation effort; before this formalisation, only a formalisation of AC-matching (which has simpler combinatorics) was reported in the proof assistant Coq [15].

Adapting first-order syntactic AC unification to the nominal setting is challenging since the new variables included in the Diophantine systems (used to generate new possible AC combinations) give rise to new AC-unification problems of the same complexity as the input problems. This paper shows that such cyclicity is not possible when only nominal AC-matching problems are considered. We present a novel nominal AC-matching algorithm adapted from the Stickel-Fages linear-Diophantine approach and prove its termination, correctness and completeness in the proof assistant PVS.

Organisation. Section 2 recalls the main concepts and notations needed in the paper. In Section 3, we present and explain the pseudocode for the algorithm specified in PVS. Section 4 discusses the main features of the formalisation, while Section 5 discusses the challenges in adapting our approach to nominal AC-unification. Finally, in Section 6, we conclude the paper and suggest possible paths for future work. This paper is an extended version of a paper submitted initially to CICM 2023.

2 Background

2.1 Nominal Terms, Permutations and Substitutions

Assume disjoint countable sets of atoms $\mathbb{A} = \{a, b, c, \ldots\}$ and of variables $\mathbb{X} = \{X, Y, Z, \ldots\}$, and a signature Σ of function symbols which contains associative-commutative function symbols. A permutation π is a bijection of the form π : $\mathbb{A} \to \mathbb{A}$ such that the domain of π (i.e., the set of atoms modified by π) is finite. Permutations are usually represented as a list of swappings, where the swapping $(a\ b)$ exchanges atoms a and b and fixes all the other atoms. Therefore, a permutation is represented as $\pi = (a_1\ b_1) :: \ldots :: (a_n\ b_n) :: nil$. The inverse of this permutation, denoted by π^{-1} , can be computed simply by reversing the list. The identity permutation is denoted by id.

Definition 1 (Nominal Terms \checkmark). The set $\mathcal{T}(\Sigma, \mathbb{A}, \mathbb{X})$ of nominal terms is generated according to the grammar:

$$s,t ::= a \mid \pi \cdot X \mid \langle \rangle \mid [a]t \mid \langle s,t \rangle \mid f t \mid f^{AC}t \tag{1}$$

where $\langle \rangle$ is the unit, a is an atom term, $\pi \cdot X$ is a moderated variable or suspension (the permutation π is suspended on the variable X), [a]t is an abstraction (a term with the atom a abstracted), $\langle s, t \rangle$ is a pair, f t is a function application and f^{AC} t is an associative-commutative function application.

Remark 1. We represent moderated variables of the form $id \cdot X$ simply as X. We follow Gabbay's name convention, which says that atoms differ in their names. Therefore, if we consider atoms a and b, it is redundant to say $a \neq b$.

Definition 2 (Well-formed Terms \square). We say that a term t is well-formed if t is not a pair and every AC-function application that is a subterm of t has at least two arguments.

As was done in [7], we have restricted the terms that our algorithm receives to well-formed terms to ease our formalisation (more details in Appendix E). Excluding pairs is a natural decision since they are used to encode a list of arguments to a function.

Definition 3 (Permutation Action). The action of permutations on atoms \Box is defined recursively: $nil \cdot c = c$ and $((a \ b) :: \pi) \cdot c = a$, if $\pi \cdot c = b$; $((a \ b) :: \pi) \cdot c = b$, if $\pi \cdot c = a$; $((a \ b) :: \pi) \cdot c = \pi \cdot c$ otherwise. The action of permutations on terms \Box is defined recursively:

$$\begin{array}{ll} \pi \cdot \langle \rangle = \langle \rangle & \pi \cdot (\pi' \cdot X) = (\pi :: \pi') \cdot X & \pi \cdot [a]t = [\pi \cdot a]\pi \cdot t \\ \pi \cdot \langle s, t \rangle = \langle \pi \cdot s, \pi \cdot t \rangle & \pi \cdot f \ t = f \ \pi \cdot t & \pi \cdot f^{AC}t = f^{AC}\pi \cdot t \end{array}$$

Notation 1 When convenient, we may mention that a function symbol f is an AC-function symbol, omit the superscript and write simply f instead of f^{AC} .

A substitution σ is a function from variables to terms, such that $\sigma X \neq id \cdot X$ only for a finite set of variables, called the domain of σ and denoted as $dom(\sigma)$. The image of σ is then defined as $im(\sigma) = \{\sigma X \mid X \in dom(\sigma)\}$. We denote the identity substitution by id. From now on, when composing substitution σ with δ we may omit the composition symbol and write $\sigma\delta$ instead of $\sigma\circ\delta$.

A well-formed substitution \checkmark only instantiates variables to well-formed terms. In the proofs of soundness and completeness of the algorithm, we restrict ourselves to well-formed substitutions. Let V be a set of variables. If $dom(\sigma) \subseteq V$ and $Vars(im(\sigma)) \subseteq V$ we write $\sigma \subseteq V$. In our PVS code, substitutions are represented by a list, where each entry of the list is called a *nuclear substitution* and is of the form $\{X \to t\}$.

Definition 4 (Nuclear substitution action on terms \Box). A nuclear substitution $\{X \to t\}$ acts over a term by induction as shown below:

$$\{X \to t\} \langle \rangle = \langle \rangle$$

$$\{X \to t\} \langle s_1, s_2 \rangle =$$

$$\{X \to t\} ([a]s) = [a](\{X \to t\}s)$$

$$\{X \to t\} (f s) = f (\{X \to t\}s)$$

$$\{X \to t\} \pi \cdot Y = \begin{cases} \pi \cdot Y \text{ if } X \neq Y \\ \pi \cdot t \text{ otherwise} \end{cases}$$

$$\{X \to t\} a = a$$

$$\{X \to t\} (f^{AC} s) = f^{AC} (\{X \to t\}s)$$

Definition 5 (Substitution acting on terms \Box **).** Since a substitution σ is a list of nuclear substitutions, the action of a substitution is defined as:

- NIL t=t, where NIL is the null list, used to represent the identity substitution.
- CONS($\{X \to s\}, \ \sigma$) $t = \{X \to s\}(\sigma t)$.

Remark 2. In the definition of action of substitutions the nuclear substitution in the head of the list is applied last. This lets us, given substitutions σ and δ , obtain the substitution $\sigma \circ \delta$ in our code simply as APPEND (σ, δ) .

2.2 Freshness and α -equality

Freshness and α -equality are two valuable notions in nominal theory and are represented by the predicates # and \approx_{α} . Intuitively, a#t means that if a occurs in t then it does so under an abstractor [a], and $s \approx_{\alpha} t$ means that s and t are α -equivalent, i.e., they are equal modulo the renaming of bound atoms. These concepts are given in Definitions 6 and 7.

Definition 6 (Freshness \Box). A freshness context ∇ is a set of constraints of the form a#X. We denote contexts by letters $\Delta, \Gamma, \nabla, \ldots$ An atom a is said to be fresh on t under a context ∇ , denoted by $\nabla \vdash a\#t$, if it is possible to build a proof using the rules:

$$\frac{\nabla \vdash a\#\langle \rangle}{\nabla \vdash a\#\langle \rangle} (\#\langle \rangle) \qquad \frac{\nabla \vdash a\#b}{\nabla \vdash a\#b} (\#atom) \qquad \frac{(\pi^{-1} \cdot a\#X) \in \nabla}{\nabla \vdash a\#\pi \cdot X} (\#X)$$

$$\frac{\nabla \vdash a\#[a]t}{\nabla \vdash a\#[b]t} (\#[a]b) \qquad \frac{\nabla \vdash a\#t}{\nabla \vdash a\#\langle s, t \rangle} (\#pair)$$

$$\frac{\nabla \vdash a\#t}{\nabla \vdash a\#f} (\#app) \qquad \frac{\nabla \vdash a\#t}{\nabla \vdash a\#f^{AC} t} (\#AC)$$

Definition 7 (α -equality with AC operators \square). Let f be an AC function symbol, $S_n(ft)$ be an operator that selects the nth argument (considering the flattened form) of ft and $D_n(ft)$ be an operator that deletes the nth argument (considering the flattened form) of ft. Then, the rule of α -equality for an AC-function application, and for some i and j, is:

$$\frac{\Delta \vdash S_i(f^{AC}s) \approx_{\alpha} S_j(f^{AC}t) \quad \Delta \vdash D_i(f^{AC}s) \approx_{\alpha} D_j(f^{AC}t)}{\Delta \vdash f^{AC}s \approx_{\alpha} f^{AC}t} (\approx_{\alpha} AC)$$

Two terms t and s are said to be α -equivalent under the freshness context Δ ($\Delta \vdash t \approx_{\alpha} s$) if it is possible to build a proof using rule ($\approx_{\alpha} AC$) and the rules:

$$\frac{\Delta \vdash \langle \rangle \approx_{\alpha} \langle \rangle}{\Delta \vdash a \approx_{\alpha} a} (\approx_{\alpha} atom)$$

$$\frac{\Delta \vdash s \approx_{\alpha} t}{\Delta \vdash f s \approx_{\alpha} f t} (\approx_{\alpha} app)$$

$$\frac{\Delta \vdash s \approx_{\alpha} t}{\Delta \vdash [a]s \approx_{\alpha} [a]t} (\approx_{\alpha} [a]a)$$

$$\frac{\Delta \vdash s \approx_{\alpha} (a b) \cdot t, \ \Delta \vdash a \# t}{\Delta \vdash [a]s \approx_{\alpha} [b]t} (\approx_{\alpha} [a]b)$$

$$\frac{\Delta \vdash s \approx_{\alpha} (a b) \cdot t, \ \Delta \vdash a \# t}{\Delta \vdash [a]s \approx_{\alpha} [b]t} (\approx_{\alpha} [a]b)$$

$$\frac{\Delta \vdash s \approx_{\alpha} (a b) \cdot t, \ \Delta \vdash a \# t}{\Delta \vdash [a]s \approx_{\alpha} [b]t} (\approx_{\alpha} [a]b)$$

$$\frac{\Delta \vdash s \approx_{\alpha} t}{\Delta \vdash [a]s \approx_{\alpha} [a]t} (\approx_{\alpha} var)$$

$$\frac{\Delta \vdash s \approx_{\alpha} t}{\Delta \vdash [a]s \approx_{\alpha} [a]t} (\approx_{\alpha} var)$$

$$\frac{\Delta \vdash s \approx_{\alpha} t}{\Delta \vdash [a]s \approx_{\alpha} [a]t} (\approx_{\alpha} var)$$

$$\frac{\Delta \vdash s \approx_{\alpha} t}{\Delta \vdash [a]s \approx_{\alpha} [a]t} (\approx_{\alpha} var)$$

$$\frac{\Delta \vdash s \approx_{\alpha} t}{\Delta \vdash [a]s \approx_{\alpha} [a]t} (\approx_{\alpha} var)$$

$$\frac{\Delta \vdash s \approx_{\alpha} t}{\Delta \vdash [a]s \approx_{\alpha} [a]t} (\approx_{\alpha} var)$$

Notation 2 We define the difference set between two permutations π and π' as $ds(\pi, \pi') = \{a \in \mathcal{A} | \pi \cdot a \neq \pi' \cdot a\}$. By extension, $ds(\pi, \pi') \# X$ is the set containing every constraint of the form a # X for $a \in ds(\pi, \pi')$.

2.3 Solution to Quintuples and Additional Notation

For the proofs of soundness and completeness of the algorithm, we need the notion of a solution to a quintuple (Definition 9). This definition depends on a parameter \mathcal{X} , a set of "protected variables", i.e., variables that cannot be instantiated.

In order to define a nominal AC-matching solution (Definition 11) and to check whether a nominal AC-equality is valid (Definition 12), we first define a general notion of unification problem (Definition 8) and a solution to a unification problem (Definition 10). Then, the definitions for nominal AC-matching and AC-equality checking are obtained immediately from the corresponding definitions of unification by correctly setting the parameter \mathcal{X} .

Definition 8 (Unification Problem). A unification problem is a triple (∇, P, \mathcal{X}) where ∇ is a freshness context; P is a finite set of equations and freshness constraints of the form $t \approx_7 s$ and $a \#_7 t$, respectively; and \mathcal{X} is a set of variables.

Notation 3 (Vars) We denote the set of variables of a term t by Vars(t) \subseteq . Let P be a finite set of equational constraints. The set of variables in P is denoted as Vars(P) \subseteq and the set of variables in the terms that are in the right-hand side of P is denoted as Vars(rhs(P)). The set of variables in $t \approx_? s$ is denoted as Vars(t,s) \subseteq . Finally, if Γ is a context then we denote by $Vars(\Gamma)$ \subseteq the set $\{X \mid a\#X \in \Gamma, \text{ for some atom } a\}$.

Notation 4 ($lhs(P) \ \ \ \ \ \$ and $rhs(P) \ \ \ \ \ \$ be a finite set of equational constraints of the form $t \approx_? s$. Do we define the left-hand side of P (denoted as lhs(P)) as the set of left-hand terms of the equational constraints in P, i.e., $lhs(P) = \{t \mid t \approx_? s \in P\}$. The right-hand side of P (denoted as rhs(P)) is defined similarly.

Notation 5 Let ∇ and ∇' be freshness contexts and σ and σ' substitutions. We need the following notation to define a solution to a quintuple:

 $-\nabla' \vdash \sigma \nabla$ denotes that $\nabla' \vdash a \# \sigma X$ holds for each $(a \# X) \in \nabla$.

 $-\nabla \vdash \sigma \approx_V \sigma'$ denotes that $\nabla \vdash \sigma X \approx_{\alpha} \sigma' X$ for all X in V. When V is the set of all variables X, we write $\nabla \vdash \sigma \approx \sigma'$.

- 1. $\Delta \vdash \delta \Gamma$. 4. there exists λ such that
- 2. if $a\#_? t \in P$ then $\Delta \vdash a\#\delta t$. $\Delta \vdash \lambda \sigma \approx_V \delta$.
- 3. if $t \approx_? s \in P$ then $\Delta \vdash \delta t \approx_\alpha \delta s$. 5. $dom(\delta) \cap \mathcal{X} = \emptyset$.

Remark 3. Note that if (Δ, δ) is a solution of $(\Gamma, \text{NIL}, \sigma, \mathbb{X}, \mathcal{X})$ this corresponds to the notion of (Δ, δ) being an instance of (Γ, σ) that does not instantiate variables in \mathcal{X} .

Definition 10 (Solution for an AC-unification problem). A solution for an AC-unification problem with protected variables (Γ, P, \mathcal{X}) is a solution for the associated quintuple $(\Gamma, P, id, Vars(P), \mathcal{X})$.

Definition 11 (Solution to an AC-matching problem). When $\mathcal{X} = Vars(rhs(P))$, Definition 8 corresponds to an AC-matching problem and Definition 10 corresponds to a solution to an AC-matching problem.

Definition 12 (Solution to an AC-equality problem). When $\mathcal{X} = Vars(P)$, Definition 8 corresponds to an AC-equality checking problem and Definition 10 corresponds to a solution to an AC-equality checking problem.

3 Algorithm

We present the algorithm's pseudocode instead of the actual PVS code for readability. We developed a nominal algorithm (Algorithm 1 \square) for matching terms t and s. The algorithm is recursive and needs to keep track of the current context Γ , the equational constraints we must unify P, the substitution σ computed so far, the set of variables V that are/were in the problem and the set of protected variables \mathcal{X} . Hence, it's input is a quintuple $(\Gamma, P, \sigma, V, \mathcal{X})$. The output is a list of solutions, each of the form (Γ_1, σ_1) . The freshness constraints are treated by auxiliary functions (see Section 3.1), and the equational constraints P are represented as a list in our PVS code, where each element of the list is a pair (t_i, s_i) that represents an equation $t_i \approx_? s_i$. The first call to the algorithm, in order to match t to s, is done with $P = \{t \approx_? s\}$; $\Gamma = \emptyset$ and $\sigma = id$ (because we have not computed any freshness constraint or substitution yet); V = Vars(t, s) and $\mathcal{X} = Vars(s)$.

Remark 4. In the PVS code, this means that the initial call is done with parameters P = cons((t, s), NIL), $\Gamma = NIL$, $\sigma = NIL$, V = Vars(t, s) and $\mathcal{X} = Vars(s)$.

Although extensive, Algorithm 1 is simple. It starts by analysing the list P of terms to match. If it is empty (line 2), it has finished and can return the answer computed so far, a list with a unique element: (Γ, σ) . Otherwise, the algorithm

Algorithm 1 Nominal AC-Matching Algorithm 1

```
1: procedure ACMATCH(\Gamma, P, \sigma, V, \mathcal{X})
        if nil?(P) then return cons((\Gamma, \sigma), NIL)
 2:
 3:
        else
             Let ((t, s), P_1) = \text{CHOOSEEQ}(P) in
 4:
             if t matches a and s matches a then return ACMATCH(\Gamma, P_1, \sigma, V, \mathcal{X})
 5:
 6:
 7:
             else if t matches \pi \cdot X and X \notin Vars(s) and X \notin \mathcal{X} then
                 \sigma_1 = \{X \mapsto \pi^{-1}s\}
 8:
                 (\Gamma_1, flag) = FRESHSUBS?(\sigma_1, \Gamma)
 9:
                 if flag then return ACMATCH(\Gamma_1 \cup \Gamma, \sigma_1 P_1, \sigma_1 \sigma, V, \mathcal{X})
10:
11:
                 else return NIL
12:
13:
              else if t matches \pi \cdot X and s matches \pi' \cdot X then
                 \Gamma_1 = ds(\pi, \pi') \# X \cup \Gamma
14:
                 return ACMATCH(\Gamma_1, P_1, \sigma, V, \mathcal{X})
15:
16:
             else if t matches \langle \rangle and s matches \langle \rangle then
17:
18:
                 return ACMATCH(\Gamma, P_1, \sigma, V, \mathcal{X})
19:
20:
             else if t matches f t_1 and s matches f s_1 then
21:
                 (P_2, flag) = DECOMPOSE(t_1, s_1)
22:
                 if flag then return ACMATCH(\Gamma, P_2 \cup P_1, \sigma, V, \mathcal{X})
23:
                 else return NIL
24:
25:
              else if t matches [a] t_1 and s = [a] s_1 then
26:
                 (P_2, flag) = DECOMPOSE(t_1, s_1)
                 if flag then return ACMATCH(\Gamma, P_2 \cup P_1, \sigma, V, \mathcal{X})
27:
28:
                 else return NIL
29:
             else if t matches [a] t_1 and s = [b]s_1 then
30:
31:
                 (\Gamma_1, flag1) = FRESH?(a, s_1)
32:
                 (P_2, flag2) = DECOMPOSE(t_1, (a b) \cdot s_1)
33:
                 if flag1 and flag2 then return ACMATCH(\Gamma \cup \Gamma_1, P_2 \cup P_1, \sigma, V, \mathcal{X})
34:
                 else return NIL
35:
             else if t matches f^{AC} t_1 and s matches f^{AC} s_1 then
36:
                 InputLst = ApplyACSTEP (\Gamma, cons((t, s), P_1), \sigma, V, X)
37:
38:
                 LstResults = MAP(ACMATCH, InputLst)
39:
                 return FLATTEN(LstResults)
40:
41:
             else return NIL
```

calls the auxiliary function CHOOSEEQ (line 4), which returns a pair (t,s) and a list of equational constraints P_1 such that $P = \{t \approx_? s\} \cup P_1$. Then, P is updated by simplifying $\{t \approx_? s\}$ and it does so by seeing the form of t (an atom, a moderated variable, a unit, and so on).

3.1 Functions ChooseEq and Decompose

The function CHOOSEEQ(P) \checkmark selects an equational constraint $t \approx_? s$ in P, picking the equation with the biggest size. This heuristic aims to aid us in the proof of termination (see Section 4.2).

The function DECOMPOSE C (lines 21, 26 and 32) receives two terms t and s, and if they are both pairs, it recursively tries to decompose them, returning a tuple (P, flag), where P is a list of equational constraints and flag is a boolean that is True if the decomposition was successful. If neither t nor s is a pair, the unification problem returned is just $P = \{t \approx_? s\}$ and flag = True. If one of the terms is a pair and the other is not, the function returns (NIL, False). In Algorithm 1, we call DECOMPOSE (t_1, s_1) when we encounter equations such as $ft_1 \approx_? fs_1$ to guarantee that all the terms in the unification problem remain well-formed. Although it would have been correct to simplify an equation of the form $ft_1 \approx_? fs_1$ to $t_1 \approx_? s_1$, if t_1 or s_1 were pairs, we would not respect our restriction that only well-formed terms are in the matching problem.

Example 1. Examples of the function DECOMPOSE are given below.

```
- DECOMPOSE(\langle a, \langle b, c \rangle \rangle, \langle c, \langle X, Y \rangle \rangle) = (\{a \approx_? c, b \approx_? X, c \approx_? Y\}, \mathit{True}).

- DECOMPOSE(a, Y) = (\{a \approx_? Y\}, \mathit{True}).

- DECOMPOSE(X, \langle c, d \rangle) = (NIL, \mathit{False}).
```

3.2 Handling Freshness Constraints - Functions freshSubs? and fresh?

Following the approach of [6], freshness constraints are handled separately by the auxiliary functions fresh? \Box and freshSubs? \Box . These functions were already implemented in [6], and extending them to handle AC-functions is straightforward. FreshSubs?(σ , Γ) returns the minimal context (Γ 1 in Algorithm 1) in which $a\#_{?}\sigma X$ holds, for every a#X in the context Γ . Fresh?(a, t) computes and returns the minimal context (Γ 1 in Algorithm 1) in which a is fresh for t. Both functions also return a boolean (flag in Algorithm 1), indicating if it was possible to find the aimed context.

3.3 The Function APPLYACSTEP

The function APPLYACSTEP \Box was adapted from the formalisation of first-order AC-unification (see [7]). It handles equations $t \approx_7 s$, where t and s are rooted by the same AC function symbol. This function returns a list (InputLst in line 37 of Algorithm 1) with each entry in this list corresponding to a branch

ACMATCH will explore. ACMATCH explores every branch generated by calling itself recursively on every input in InputLst (line 38 of the algorithm). The algorithm's output is a list of solutions of the form (Γ, σ) , where Γ is a context and σ is a substitution. In addition, the result of calling MAP(ACMATCH, InputLst), LstResults in line 38 of Algorithm 1, is a list of lists of solutions. Hence, LstResults is flattened and then returned.

Remark 5 (SOLVEAC and INSTANTIATESTEP). APPLYACSTEP relies on two functions: SOLVEAC and INSTANTIATESTEP , which are fully described in [7]. In synthesis, the function SOLVEAC finds the linear Diophantine equational system associated with the AC-matching equational constraint, generates the basis of solutions, and uses these solutions to generate the new AC-matching equational constraints. The function INSTANTIATESTEP instantiates the moderated variables that it can.

4 Formalisation

As is done in [7], to help us in the proofs of termination (Section 4.2), soundness (Section 4.3) and completeness (Section 4.4) we define the notion of a *nice input* (Section 4.1). More details about the PVS formalisation can be found in Appendix B.

4.1 Nice Inputs

Nice inputs are invariant under the action of the ACMATCH function with valuable properties. Notice that Item 7 of Definition 13 would need to be removed for the proofs of termination, soundness, and completeness to be used in unification.

Definition 13 (Nice input \mathcal{C}). An input $(\Gamma, P, \sigma, V, \mathcal{X})$ is said to be nice if:

```
1. \sigma is idempotent.

2. Vars(P) \cap dom(\sigma) = \emptyset.

3. \sigma \subseteq V.

4. Vars(P) \subseteq V.

5. Vars(\Gamma) \subseteq V.

6. \mathcal{X} \subseteq V.

7. Vars(rhs(P)) \subseteq \mathcal{X}.
```

Items 1 to 4 were present in the definition of nice input for the formalisation of first-order AC-unification, while items 5 to 7 were added. Item 5 of Definition 13 was expected, since we already have similar hypotheses for P and σ . Item 6 guarantees that the new variables introduced by the algorithm can be instantiated (see Appendix A).

4.2 Termination

For the lexicographic measure used in the proof of termination, we need the definition of the size of an equational constraint $t \approx_7 s$ (Definition 14).

Definition 14 (Size of an Equational Constraint \checkmark). The size of an equational constraint $t \approx_? s$ is size(t) + size(s), where the size of a term $t \checkmark$ is recursively defined as follows:

```
\begin{array}{ll} -\ size(a) = 1. \\ -\ size(\pi \cdot X) = 1. \\ -\ size(\langle\rangle) = 1. \end{array} \\ \begin{array}{ll} -\ size(\langle t_1, t_2 \rangle) = 1 + size(t_1) + size(t_2). \\ -\ size(f\ t_1) = 1 + size(t_1). \\ -\ size(f^{AC}\ t_1) = 1 + size(t_1). \\ -\ size([a]t_1) = 1 + size(t_1) \end{array}
```

Although the nominal AC-matching algorithm is based on the first-order AC-unification algorithm ([7]), the proof of termination was much easier for nominal AC-matching than for first-order AC-unification. Instead of the intricate lexicographic measure used in [7] (which came from the work of [17]), it was possible to prove that for the particular case of matching (unlike unification) all the new moderated variables introduced by SOLVEAC are instantiated by INSTANTIATESTEP.

Hence, the lexicographic measure used has as its first component the number of variables in the equational constraints P and as a second component the multiset order of the size of each equation $t \approx_? s \in P$. Although PVS does not directly implement multiset orders, this part can be emulated easily by analysing the maximum size n of all equations $t \approx_? s$ in P and the number of equations $t \approx_? s$ in P with maximal size (in this order). The algorithm selects an equation with maximal size to simplify (the heuristic selection is enforced by the function CHOOSEEQ).

Let MS(P) be the maximum size n of all equations $t \approx_? s$ in P and let NMS(P) be the number of equations $s \approx_? t$ whose size is equal to MS(P). The lexicographic measure in then lex = (|Vars(P)|, MS(P), NMS(P)). Table 1 shows which component do not increase (represented by \leq) and which components strictly decrease (represented by \leq).

Recursive Call	Vars(P)	MS(P)	NMS(P)
line 5, 15, 18, 22, 27, 33, 38 (case 1)	<u> </u>	\leq	<
line 5, 15, 18, 22, 27, 33, 38 (case 2)	<	<	

Table 1. Decrease of the components of the lexicographic measure.

4.3 Soundness

As mentioned, to match terms t and s we first call the Algorithm 1 with parameters $\Gamma = \emptyset$, $P = \{t \approx_? s\}$, $\sigma = id$, V = Vars(t,s) and $\mathcal{X} = Vars(s)$. However, since the parameters of ACMATCH change after recursive calls, the proof of soundness (Corollary 1) cannot be done directly by induction, and we must instead prove first the Theorem 1 with generic parameters Γ , P, σ , V and \mathcal{X} . Once the Theorem 1 is proved, it is also immediate to adapt the algorithm to solve nominal AC-equality checking and to prove its soundness (Corollary 2).

Theorem 1 (Soundness for Nice Inputs \mathcal{C}). Let the pair (Γ_1, σ_1) an output of ACMATCH $(\Gamma, P, \sigma, V, \mathcal{X})$ and suppose that $(\Gamma, P, \sigma, V, \mathcal{X})$ is a nice input. If (Δ, δ) is a solution to $(\Gamma_1, \text{NIL}, \sigma_1, \mathbb{X}, \mathcal{X})$ then (Δ, δ) is a solution to $(\Gamma, P, \sigma, \mathbb{X}, \mathcal{X})$.

Corollary 1 (Soundness for AC-Matching \Box). Let the pair (Γ_1, σ_1) an output of ACMATCH(\emptyset , $\{t \approx_? s\}$, id, Vars(t, s), Vars(s)). If (Δ, δ) is an instance of (Γ_1, σ_1) that does not instantiate the variables in s, then (Δ, δ) is a solution to $(\emptyset, \{t \approx_? s\}, id, \mathbb{X}, Vars(s))$.

Corollary 2 (Soundness for AC-Equality Checking $oldsymbol{C}$). Let (Γ_1, σ_1) be an output of ACMATCH(\emptyset , $\{t \approx_? s\}$, id, Vars(t, s), Vars(t, s)). If (Δ, δ) is an instance of (Γ_1, σ_1) that does not instantiate the variables in t or s, then (Δ, δ) is a solution to $(\emptyset, \{t \approx_? s\}, id, \mathbb{X}, Vars(t, s))$.

Remark 6. An interpretation of Corollary 1 is that if (Δ, δ) is an AC-matching instance to one of the outputs of ACMATCH, then (Δ, δ) is an AC-matching solution to the original problem. Corollary 2 has a similar interpretation, replacing AC-matching with AC-equality checking.

The proof of soundness was mainly a straightforward adaptation from the proof of soundness of first-order AC-unification ([7]). The soundness of FRESH? and FRESHSUBS? was a straightforward adaptation from the work of [6] since the only case not covered in [6] (the case of AC-functions) is similar to the case of syntactic functions.

4.4 Completeness

Completeness of Algorithm 1 is given by the Corollary 3 and similarly to the soundness proof, it is derived easily after proving the Theorem 2.

Theorem 2 (Completeness for Nice Inputs $oldsymbol{C}$). Let $(\Gamma, P, \sigma, V, \mathcal{X})$ be a nice input. Suppose that (Δ, δ) is a solution to $(\Gamma, P, \sigma, \mathbb{X}, \mathcal{X})$, that $\delta \subseteq V$ and that $Vars(\Delta) \subseteq V$. Then, there exists $(\Gamma_1, \sigma_1) \in ACMATCH(\Gamma, P, \sigma, V, \mathcal{X})$ such that (Δ, δ) is an instance (restricted to the variables of V) of (Γ_1, σ_1) that does not instantiate the variables in \mathcal{X} .

Corollary 3 (Completeness for AC-Matching \square). Suppose that (Δ, δ) is a solution to $(\emptyset, \{t \approx_? s\}, id, \mathbb{X}, Vars(s))$, that $\delta \subseteq V$ and that $Vars(\Delta) \subseteq V$. Then, there exists $(\Gamma_1, \sigma_1) \in ACMATCH(\emptyset, \{t \approx_? s\}, id, V, Vars(s))$ such that (Δ, δ) is an instance (restricted to the variables of V) of (Γ_1, σ_1) that does not instantiate the variables of s.

Corollary 4 (Completeness for AC-equality Checking \checkmark). Suppose (Δ, δ) is a solution to $(\emptyset, \{t \approx_? s\}, id, \mathbb{X}, Vars(t, s))$ satisfying $\delta \subseteq V$ and $Vars(\Delta) \subseteq V$. Then, there exists $(\Gamma_1, \sigma_1) \in ACMATCH(\emptyset, \{t \approx_? s\}, id, V, Vars(t, s))$ such that (Δ, δ) is an instance (restricted to the variables of V) of (Γ_1, σ_1) that does not instantiate the variables of t or s.

Remark 7. An interpretation of Corollary 3 is that if (Δ, δ) is an AC-matching solution to the initial problem, then (Δ, δ) is an AC-matching instance of one of the outputs of ACMATCH. Corollary 4 has a similar interpretation, replacing AC-matching with AC-equality checking.

As was the case for first-order AC-unification (see [7]), the hypothesis $\delta \subseteq V$ in the proof of completeness is merely a technicality that was put in order to guarantee the new variables introduced by the algorithm in the AC-part do not clash with the variables in $dom(\delta)$ or in the terms in $im(\delta)$. This mechanism could be replaced by a different one that assures that the variables introduced by the AC-part of ACMATCH are indeed new. When going from the first-order setting to the nominal setting, we go from having a unifier δ to a pair (Δ, δ) and hence we must add the hypothesis $Vars(\Delta) \subseteq V$.

Remark 8 (High-level description of how to remove hypotheses $\delta \subseteq V$ and $Vars(\Delta) \subseteq V$). The critical step to prove a variant of Corollary 3 with V = Vars(t,s) and without the hypothesis $\delta \subseteq V$ and $Vars(\Delta) \subseteq V$ is to prove that the outputs computed when we call ACMATCH with input $(\Gamma, P, \sigma, V, \mathcal{X})$ "differ only by a renaming" from the outputs computed when we call ACMATCH with input $(\Gamma, P, \sigma, V', \mathcal{X})$. However, this cannot be proved directly by induction because if V and V' differ and ACMATCH enters in the AC-part, the new variables introduced for each input may "differ only by a renaming" and once we instantiate those variables, it may happen that the substitutions computed so far (the third component in the input quintuple) will also "differ only by a renaming". As was done for first-order AC-unification, the solution is to prove the more general statement that if the inputs $(\Gamma, P, \sigma, V, \mathcal{X})$ and $(\Gamma', P', \sigma', V', \mathcal{X}')$ "differ only by a renaming", then the output of when we call ACMATCH with the first input "differ only by a renaming" from the output of when we call ACMATCH with the second input.

5 Towards a Nominal AC-Unification Algorithm

Stickel's AC-unification algorithm relies on solving Diophantine equations where new variables are used to represent arguments of AC operators. Using the same approach to solve nominal AC-unification problems leads to non-termination in cases where the same variable occurs as an argument of an AC operator multiple times with *different* suspended permutations.

In the case of nominal AC-matching, the hypothesis $Vars(rhs(P)) \subseteq \mathcal{X}$ in the definition of nice inputs is essential to prove termination. For example, if we are working under a freshness context Γ and apply ACMATCH to the equational constraint $f(X,W) \approx_? f(\pi \cdot X, \pi \cdot Y)$, where $X \notin \mathcal{X}$ and we do not have $\Gamma \vdash X \approx_\alpha \pi \cdot X$, we obtain a problem that is essentially a renaming of our original problem (we would compute the substitution $\sigma = \{X \mapsto f(Y_1, X_1), W \mapsto W_1, Y \mapsto \pi^{-1} \cdot Y_1\}$ and be left with the equational constraint $f(X_1, W_1) \approx_? f(\pi \cdot X_1, \pi \cdot Y_1)$). More details can be found in Appendix C.

This problem does not arise in first-order AC-unification because, in the corresponding first-order problem, we would not have two different permutations (id and π in this case) suspended on the same variable (X in this case). Instead, we would have the same variable X as an argument to both terms and eliminate it. This problem does not arise in nominal AC-matching because X would be a protected variable. Hence, we would not compute the substitution $\sigma = \{X \mapsto f(Y_1, X_1), W \mapsto W_1, Y \mapsto \pi^{-1} \cdot Y_1\}$, we would instead discard this branch. In future work, we will consider the alternative approach to AC-unification proposed by Boudet, Contenjean and Devie [10,8], which was used to define AC higher-order pattern unification [9]. To our knowledge, this AC unification approach has not been formalised yet. However, it has the advantage of generating simpler Diophantine systems, which could be easier to combine with nominal unification.

6 Conclusion and Future Work

We propose the first (to the best of our knowledge) nominal AC-matching algorithm, together with proofs of its termination, soundness and completeness. All proofs were formalised in the proof assistant PVS. As a byproduct, we also obtained a formalised nominal AC-equality checking algorithm. Nominal AC-matching has applications for nominal AC-rewriting, being the first step towards a nominal AC-unification algorithm.

Our formalisation extends the formalisation of first-order AC-unification by Ayala-Rincón et al. [7] to nominal terms and uses the functions that deal with freshness constraints from [6], extending them to deal with AC-function symbols. Furthermore, by adding a parameter \mathcal{X} for protected variables, it enables both AC-matching and AC-equality checking, according to whether \mathcal{X} is the set of variables in the right-hand side of the problem or the set of variables in the problem. The .pvs files have a combined size of 290 KB and contain the specification of functions and the statements of the theorems. The .prf files contain the proofs of the theorems and have a combined size of 22 MB.

Future work will explore ways to define a nominal AC-unification algorithm, avoiding the loop described in Section 5. We will consider alternative AC-unification algorithms as a starting point [10,9] and explore the connection between higher-order pattern unification and nominal unification (e.g., [13,22]).

A nominal AC-unification algorithm would have applications in logic programming languages that employ the nominal paradigm, such as α -Prolog. A second possible future work path is to use this formalisation to formalise a more efficient nominal AC-matching algorithm. Finally, a third future work path would be formalising matching/unification algorithms for different equational theories.

Acknowledgenments. Partially supported by the Austrian Science Fund (FWF), Project P 35530, Brazilian FAP-DF Project DE 00193.00001175/2021-11 and, Brazilian CNPq Project Universal 409003/2021-2. First author partially funded by a CNPq productivity research grant 313290/2021-0.

References

- Ayala-Rincón, M., de Carvalho Segundo, W., Fernández, M., Nantes-Sobrinho, D.: Nominal C-Unification. In: Logic-Based Program Synthesis and Transformation - 27th International Symposium, LOPSTR, Revised Selected Papers. LNCS, vol. 10855, pp. 235–251. Springer (2017). https://doi.org/10.1007/978-3-319-94460-9 14
- Ayala-Rincón, M., de Carvalho Segundo, W., Fernández, M., Nantes-Sobrinho, D.: On Solving Nominal Fixpoint Equations. In: Frontiers of Combining Systems -11th International Symposium, FroCoS. LNCS, vol. 10483, pp. 209–226. Springer (2017). https://doi.org/10.1007/978-3-319-66167-4 12
- 3. Ayala-Rincón, M., de Carvalho Segundo, W., Fernández, M., Nantes-Sobrinho, D., Oliveira, A.C.R.: A formalisation of nominal α -equivalence with A, C, and AC function symbols. Theor. Comput. Sci. **781**, 3–23 (2019). https://doi.org/10.1016/j.tcs.2019.02.020
- Ayala-Rincón, M., de Carvalho Segundo, W., Fernández, M., Silva, G.F., Nantes-Sobrinho, D.: Formalising Nominal C-Unification Generalised with Protected Variables. Math. Struct. Comput. Sci. 31(3), 286–311 (2021). https://doi.org/10.1017/S0960129521000050
- Ayala-Rincón, M., Fernández, M., Nantes-Sobrinho, D.: On Nominal Syntax and Permutation Fixed Points. Log. Methods Comput. Sci. 16(1) (2020). https://doi.org/10.23638/LMCS-16(1:19)2020
- Ayala-Rincón, M., Fernández, M., Oliveira, A.C.R.: Completeness in PVS of a Nominal Unification Algorithm. In: Proceedings of the Tenth Workshop on Logical and Semantic Frameworks, with Applications, LSFA. ENTCS, vol. 323, pp. 57–74. Elsevier (2015). https://doi.org/10.1016/j.entcs.2016.06.005
- Ayala-Rincón, M., Fernández, M., Silva, G.F., Sobrinho, D.N.: A Certified Algorithm for AC-Unification. In: 7th International Conference on Formal Structures for Computation and Deduction, FSCD. LIPIcs, vol. 228, pp. 8:1–8:21. Schloss Dagstuhl Leibniz-Zentrum für Informatik (2022). https://doi.org/10.4230/LIPIcs.FSCD.2022.8
- Boudet, A.: Competing for the AC-Unification Race. J. of Autom. Reasoning 11(2), 185–212 (1993). https://doi.org/10.1007/BF00881905
- 9. Boudet, A., Contejean, E.: AC-Unification of Higher-Order Patterns. In: Third International Conference on Principles and Practice of Constraint Programming CP97. LNCS, vol. 1330, pp. 267–281. Springer (1997). https://doi.org/10.1007/BFb0017445
- Boudet, A., Contejean, E., Devie, H.: A New AC Unification Algorithm with an Algorithm for Solving Systems of Diophantine Equations. In: Proceedings of the Fifth Annual Symposium on Logic in Computer Science, LICS. pp. 289–299. IEEE Computer Society (1990). https://doi.org/10.1109/LICS.1990.113755
- 11. Calvès, C.F., Fernández, M.: Matching and Alpha-Equivalence Check for Nominal Terms. J. of Computer and System Sciences **76**(5), 283–301 (2010). https://doi.org/http://dx.doi.org/10.1016/j.jcss.2009.10.003
- 12. Calvès, C., Fernández, M.: A polynomial nominal unification algorithm. Theor. Comput. Sci. 403(2-3), 285–306 (2008). https://doi.org/10.1016/j.tcs.2008.05.012
- 13. Cheney, J.: Relating nominal and higher-order pattern unification. In: Proceedings of the 19th international workshop on Unification, UNIF. pp. 104–119 (2005)
- Clausen, M., Fortenbacher, A.: Efficient Solution of Linear Diophantine Equations.
 J. of Sym. Computation 8(1-2), 201–216 (1989). https://doi.org/10.1016/S0747-7171(89)80025-2

- Contejean, E.: A Certified AC Matching Algorithm. In: Rewriting Techniques and Applications, 15th International Conference, RTA. LNCS, vol. 3091, pp. 70–84. Springer (2004). https://doi.org/10.1007/978-3-540-25979-4
- 16. Eker, S.: Associative-Commutative Rewriting on Large Terms. In: Rewriting Techniques and Applications, 14th International Conference, RTA. LNCS, vol. 2706, pp. 14–29. Springer (2003). https://doi.org/10.1007/3-540-44881-0 3
- Fages, F.: Associative-Commutative Unification. In: 7th International Conference on Automated Deduction CADE. LNCS, vol. 170, pp. 194–208. Springer (1984). https://doi.org/10.1007/978-0-387-34768-4 12
- 18. Fages, F.: Associative-Commutative Unification. J. of Sym. Computation **3**(3), 257–275 (1987). https://doi.org/10.1016/S0747-7171(87)80004-4
- 19. Fernández, M., Gabbay, M.J.: Nominal rewriting. Information and Computation **205**(6), 917–965 (2007). https://doi.org/10.1016/j.ic.2006.12.002
- Gabbay, M.J., Pitts, A.M.: A New Approach to Abstract Syntax with Variable Binding. Formal Aspects of Computing 13(3), 341–363 (Jul 2002). https://doi.org/10.1007/s001650200016
- Levy, J., Villaret, M.: An Efficient Nominal Unification Algorithm. In: Proceedings of the 21st International Conference on Rewriting Techniques and Applications, RTA. LIPIcs, vol. 6, pp. 209–226. Schloss Dagstuhl - Leibniz-Zentrum für Informatik (2010). https://doi.org/10.4230/LIPIcs.RTA.2010.209
- Levy, J., Villaret, M.: Nominal Unification from a Higher-Order Perspective. ACM Trans. Comput. Log. 13(2), 10:1–10:31 (2012). https://doi.org/10.1145/2159531.2159532
- 23. Pitts, A.M.: Nominal Sets: Names and Symmetry in Computer Science. Cambridge University Press (2013)
- Pottier, F.: An Overview of CαML. In: Benton, N., Leroy, X. (eds.) Proceedings of the ACM-SIGPLAN Workshop on ML, ML 2005, Tallinn, Estonia, September 29, 2005. Electronic Notes in Theoretical Computer Science, vol. 148, pp. 27–52. Elsevier (2005). https://doi.org/10.1016/j.entcs.2005.11.039
- Schmidt-Schauß, M., Kutsia, T., Levy, J., Villaret, M.: Nominal Unification of Higher Order Expressions with Recursive Let. In: Logic-Based Program Synthesis and Transformation 26th International Symposium, LOPSTR, Revised Selected Papers. LNCS, vol. 10184, pp. 328–344. Springer (2016). https://doi.org/10.1007/978-3-319-63139-4 19
- Schmidt-Schauß, M., Kutsia, T., Levy, J., Villaret, M., Kutz, Y.D.K.: Nominal Unification and Matching of Higher Order Expressions with Recursive Let. Fundam. Informaticae 185(3), 247–283 (2022). https://doi.org/10.3233/FI-222110
- 27. Shinwell, M.R., Pitts, A.M., Gabbay, M.: FreshML: programming with binders made simple. In: Runciman, C., Shivers, O. (eds.) Proceedings of the Eighth ACM SIGPLAN International Conference on Functional Programming, ICFP 2003, Uppsala, Sweden, August 25-29, 2003. pp. 263–274. ACM (2003). https://doi.org/10.1145/944705.944729
- 28. Stickel, M.E.: A Complete Unification Algorithm for Associative-Commutative Functions. In: Advance Papers of the Fourth International Joint Conference on Artificial Intelligence, IJCAI. pp. 71–76 (1975), http://ijcai.org/Proceedings/75/Papers/011.pdf
- 29. Stickel, M.E.: A Unification Algorithm for Associative-Commutative Functions. J. of the ACM $\bf 28(3)$, $\bf 423-434$ (1981). https://doi.org/10.1145/322261.322262
- 30. Urban, C.: Nominal Techniques in Isabelle/HOL. J. Autom. Reason. **40**(4), 327–356 (2008). https://doi.org/10.1007/s10817-008-9097-2

- 31. Urban, C.: Nominal Unification Revisited. In: Proceedings 24th International Workshop on Unification, UNIF. EPTCS, vol. 42, pp. 1–11 (2010). https://doi.org/10.4204/EPTCS.42.1
- 32. Urban, C., Pitts, A.M., Gabbay, M.: Nominal unification. Theor. Comput. Sci. **323**(1-3), 473–497 (2004). https://doi.org/10.1016/j.tcs.2004.06.016

A An Example of First-Order AC-Unification and How to Adapt it to Nominal AC-Matching

Notation 6 (Flattened form of AC-functions) Let f be an AC-function symbol. When convenient, we may denote in this paper an AC-function in flattened form. For instance, the term $f\langle\langle a,b\rangle,\langle c,d\rangle\rangle$ may be denoted simply as f(a,b,c,d). In our formalisation, when we manipulate an AC-function term t we are more interested in its arguments than in how they were encoded using pairs.

Terms such as $f^{AC}\langle\langle a,b\rangle, f^{AC}\langle c,d\rangle\rangle$, $f^{AC}\langle f^{AC}\langle\rangle, f^{AC}\langle c,d\rangle\rangle$, and $f^{AC}\langle\langle\rangle, f^{AC}\langle\langle c,d\rangle\rangle$ are denoted respectively by f(a,b,c,d), $f(\langle\rangle,c,d)$ and $f(\langle\rangle,c,d)$.

A.1 An Example

We give a higher-level example (taken from the very accessible [29] and also present in [7]) of how we would solve the first-order AC-unification problem $\{f(X,X,Y,a,b,c)\approx_? f(b,b,b,c,Z)\}$, where f is an AC-function symbol. In a high-level view, this technique converts an AC-unification problem into a linear Diophantine equation and uses a basis of solutions of the Diophantine equation to get a complete set of AC-unifiers to our original problem.

The first step is to eliminate common arguments in the terms that we are trying to unify. The problem is now $\{f(X,X,Y,a)\approx_? f(b,b,Z)\}$. The second step is to associate our unification problem with a linear Diophantine equation, where each argument of our terms corresponds to one variable in the equation (this process is called variable abstraction) and the coefficient of this variable in the equation is the number of occurrences of the argument. In our case, the linear Diophantine equation obtained is: $2X_1 + X_2 + X_3 = 2Y_1 + Y_2$ (variable X_1 was associated with argument X, variable X_2 with the argument Y and so on; the coefficient of variable X_1 is two, since argument X occurs twice in f(X,X,Y,a) and so on).

Tabl	e 2.	Solutions	for	$_{ m the}$	equation	$2X_1$	+	$X_2 +$	$X_3 =$	$= 2Y_1 +$	Y_2 .
------	------	-----------	-----	-------------	----------	--------	---	---------	---------	------------	---------

X	1	X_2	X_3	$ Y_1 $	Y_2	New Variables.
-0)	0	1	0	1	Z_1
C)	1	0	0	1	Z_2
C)	0	2	1	0	Z_3
0)	1	1	1	0	Z_4
0)	2	0	1	0	Z_5
1		0	0	0	2	Z_6
1		0	0	1	0	Z_7

The third step is to generate a basis of solutions to the equation and associate a new variable (the Z_i s) to each solution. The result is shown on Table 2. As we will soon see, the unification problem $\{f(X, X, Y, a) \approx_? f(b, b, Z)\}$ may branch

into (possibly) many unification problems and the new variables Z_i s will be the building blocks for the right-hand side of these unification problems. Observing Table 2 we relate the "old variables" (X_i s and Y_i s) with the "new variables" (Z_i s):

$$X_{1} = Z_{6} + Z_{7}$$

$$X_{2} = Z_{2} + Z_{4} + 2Z_{5}$$

$$X_{3} = Z_{1} + 2Z_{3} + Z_{4}$$

$$Y_{1} = Z_{3} + Z_{4} + Z_{5} + Z_{7}$$

$$Y_{2} = Z_{1} + Z_{2} + 2Z_{6}.$$
(2)

In order to explore all possible solutions, we must consider whether we will include or not each solution on our basis. Since seven solutions compose our basis (one for each variable Z_i), this means that a priori there are 2^7 cases to consider. Considering that including a solution of our basis means setting the corresponding variable Z_i to 1 and not including it means setting it to 0, we must respect the constraint that no original variables $(X_1, X_2, X_3, Y_1, Y_2)$ receive 0. Eliminating the cases that do not respect this constraint, we are left with 69 cases.

For example, if we decide to include only the solutions represented by the variables Z_1 , Z_4 and Z_6 , the corresponding unification problem, according to the Equations (2), becomes:

$$P = \{X_1 \approx_? Z_6, X_2 \approx_? Z_4, X_3 \approx_? f(Z_1, Z_4), Y_1 \approx_? Z_4, Y_2 \approx_? f(Z_1, Z_6, Z_6)\}.$$
(3)

We can also drop the cases where a variable that does not represent a variable term is paired with an AC-function application. For instance, the unification problem P should be discarded, since the variable X_3 represents the constant a, and we cannot unify a with $f(Z_1, Z_4)$. This constraint eliminates 63 of the 69 potential unifiers.

Finally we replace the variables X_1, X_2, X_3, Y_1, Y_2 by the original arguments they substituted and proceed with the unification. Some unification problems that we will explore will be unsolvable and discarded later, as:

$$\{X \approx_? Z_6, Y \approx_? Z_4, a \approx_? Z_4, b \approx_? Z_4, Z \approx_? f(Z_6, Z_6)\}$$

(we cannot unify both a with Z_4 and b with Z_4 simultaneously). In the end, the solutions computed will be:

$$\sigma_{1} = \{Y \to f(b,b), Z \to f(a,X,X)\},
\sigma_{2} = \{Y \to f(Z_{2},b,b), Z \to f(a,Z_{2},X,X)\},
\sigma_{3} = \{X \to b, Z \to f(a,Y)\},
\sigma_{4} = \{X \to f(Z_{6},b), Z \to f(a,Y,Z_{6},Z_{6})\}.$$
(4)

Remark 9. When using the technique described in this section to unify f(X, X, Y, a, b, c) with f(b, b, b, c, Z), we obtained unification problems that only contain

the variables X_1 , X_2 , X_3 , Y_1 , Y_2 or AC-functions whose arguments are all variables (for instance P in Equation 3). However, this does not mean that our technique cannot be applied to general AC-unification problems, since we eventually replace the variables X_1 , X_2 , X_3 , Y_1 , Y_2 by their corresponding arguments (X, Y, a, b, Z) respectively) and proceed with unification.

A.2 Modifications to Adapt the Algorithm to the Nominal Setting

The example describes the process of trying to unify two terms $t \equiv f(t_1, \ldots, t_m)$ and $s \equiv f(s_1, \ldots, s_n)$, where f is an AC-function symbol. Four modifications were necessary to adapt this process to the nominal setting.

The first is related to eliminating common arguments: we do not eliminate arguments t_i and s_j of t and s if they are equal modulo AC, we eliminate them if they are α -equivalent (modulo AC) under the context Γ that we are working with, i.e., if $\Gamma \vdash t_i \approx_{\alpha} s_j$. If we have as hypothesis that (Δ, δ) is the solution to the quintuple we are working with (see Definition 9), the correctness of this step boils down to proving that from $\Gamma \vdash t_i \approx_{\alpha} s_j$ we have $\Delta \vdash \delta t_i \approx_{\alpha} \delta s_j$. This is possible to prove by using the fact that $\Delta \vdash \delta \Gamma$ (item 1 of Definition 9).

The second change is related with the new variables (Z_i s in the example of Section A.1) introduced and the fact that in the nominal setting a moderated variable $\pi \cdot X$ always has a permutation π suspended on the variable X. What should be the permutation π suspended on the new variables? Since the ultimate goal of these new variables is to outline the combinatory between the arguments of t and the arguments of s, we put the identity permutation suspended on the new variables. For instance, in Example of Section A.1 we would have the moderated variables $id \cdot Z_1, \ldots, id \cdot Z_7$, which we would write (see Remark 1) simply as Z_1, \ldots, Z_7 .

In Example of Section A.1, we have variables X_1, X_2, X_3, Y_1, Y_2 to represent respectively the arguments X, Y, a, b, Z and we say that when generating the new unification problems we can discard the ones "where a variable that does not represent a variable term is paired with an AC-function application". Here, we can also discard unification problems where a moderated variable $\pi \cdot X$, with $X \in \mathcal{X}$, is paired with an AC-function application. This is the third change to adapt to the nominal setting.

Finally, we must guarantee that the new variables Z_i s introduced by the algorithm can be instantiated. Since those new variables are not in the set V, we ensure that by putting the restriction that $\mathcal{X} \subseteq V$ in the definition of nice inputs (Definition 13).

B More Information About the PVS Formalisation

The functions coded in PVS and the statement of the theorems can be found in files .pvs, while the proofs of the theorems can be found in the .prf files. Figure 1 shows the dependency diagram for the PVS theories of the formalisation. An

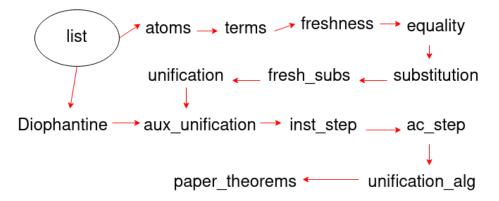


Fig. 1. Dependency Diagram for PVS Theories.

arrow going from theoryA to theoryB means that theoryA imports definitions and lemmas from theoryB.

A short description of each of these theories is show below:

- paper_theorem highlights the most important theorems that appear in this paper.
- unification_alg contains the function ACMATCH and the lemmas that directly support the proofs of soundness and completeness.
- ac_step contains function APPLYACSTEP and lemmas about its properties.
- inst_step contains function INSTANTIATESTEP and related lemmas.
- aux_unification contains functions SOLVEAC (with lemmas about its properties) and the main functions called by SOLVEAC (with lemmas about its properties).
- diophantine definitions and properties about solving linear Diophantine equations.
- unification definition of solution to a quintuple and lemmas about unification.
- fresh_subs definition and properties of FRESHSUBS?.
- substitution definition and properties about substitutions.
- equality definition and properties about nominal AC-equality checking.
- freshness definition and properties about freshness. Contain function FRESH?
- terms definition and properties about terms.
- atoms definition and properties of permutations and their actions on atoms.
- list this is a set of parametric theories that define generic functions that operate on lists, not strictly connected to unification.

When specifying theorems and functions, PVS may generate proof obligations that must be proved by the user. These proof obligations are called Type Correctness Conditions (TCCs) and the PVS system includes several predefined

proof strategies that automatically discharge most of the TCCs. In our specification, most TCCs were related to the termination of functions and PVS was able to prove almost all of them automatically. The number of theorems and TCCs proved for each theory, along with the approximate size of each theory and their percentage of the total size is shown in Table 3.

Theory	Theorems	TCCs	Size (.pvs)	Size (.prf)	Size (%)
paper_theorems	6	4	2.8 kB	0.02 MB	< 1%
unification_alg	11	19	6.9 kB	2.1 MB	9%
ac_step	45	11	15.8 kB	1.6 MB	7%
inst_step	75	17	20.3 kB	2 MB	9%
aux_unification	140	52	44.9 kB	6.9 MB	30%
Diophantine	77	44	$23.5~\mathrm{kB}$	1 MB	4%
unification	119	13	28.0 kB	1.7 MB	8%
fresh_subs	37	5	10.9 kB	0.6 MB	3%
substitution	166	34	30.1 kB	2.5 MB	11%
equality	83	20	15.1 kB	1.6 MB	7%
freshness	15	10	4.5 kB	0.1 MB	< 1%
terms	147	53	29.1 kB	1.1 MB	5 %
atoms	14	3	3.7 kB	0.03 MB	< 1 %
list	265	113	54.9 kB	1.4 MB	6 %
Total	1200	398	290.5 kB	22 6MB	100%

Table 3. Main Information on the Theories of Our Formalisation.

Remark 10. The theory atoms has its definitions and lemmas in the file atoms.pvs and the proofs of the lemmas in the file atoms.prf. The same happens for all the theories mentioned in this diagram, except list. In Figure 1, list represents a set of parametric theories that define generic functions (not strictly connected to matching) that operate on lists. The theories in list are list_nat_theory, list_theory, list_theory, map_theory and more_list_theory_props. However, since the specifics of each theory in list is not significant to our formalisation, we grouped them together in our diagram.

C The Loop in Nominal AC-Unification

Notation 7 From now on, when denoting a suspended variable $\pi \cdot X$ we may omit the \cdot symbol and write simply πX .

Suppose that we are working under an empty context (i.e. $\Gamma = \emptyset$) and want to solve the equational constraint $f(X,W) \approx_? f(\pi X,\pi Y)$, with $\mathcal{X} = \emptyset$. The linear Diophantine equation associated with this problem $U_1 + U_2 = V_1 + V_2$, where variable U_1 is associated with argument X, variable U_2 is associated with argument W, variable V_1 is associated with argument T and variable T is associated with argument T and T are T and T are T and T and T are T and T are T are T and T are T are T and T are T and T are T are T and T are T are T and T are T and T are T are T and T are T and T are T and T are T and T are T are T and T are T are T and T are T are T are T and T are T and T are T are T and T are T are T and T are T and T are T are T and T are T and T are T are T and T are T and T are T are T and T are T and T

Table 4. Solutions for the Equation $U_1 + U_2 = V_1 + V_2$

U_1	U_2	V_1	V_2	$ U_1+U_2 $	$V_1 + V_2$	New Variables
0	1	0	1	1	1	Z_1
0	1	1	0	1	1	W_1
1	0	0	1	1	1	Y_1
1	0	1	0	1	1	X_1

associated with argument πY . A basis of solutions to this linear Diophantine equation is shown in Table C.

We choose the names of the new variables to be Z_1 , W_1 , Y_1 and X_1 deliberately to make the loop in nominal AC-unification clearer. Finally, we will branch into new equational constraints, using Table C to construct them. The algorithm bifurcates into 7 branches, shown below along with their corresponding equational constraints:

$$\begin{aligned} branch_1 &= \{X \approx_? X_1, W \approx_? Z_1, \pi X \approx_? X_1, \pi Y \approx_? Z_1\} \\ branch_2 &= \{X \approx_? Y_1, W \approx_? W_1, \pi X \approx_? W_1, \pi Y \approx_? Y_1\} \\ branch_3 &= \{X \approx_? Y_1 + X_1, W \approx_? W_1, \pi X \approx_? W_1 + X_1, \pi Y \approx_? Y_1\} \\ branch_4 &= \{X \approx_? Y_1 + X_1, W \approx_? Z_1, \pi X \approx_? X_1, \pi Y \approx_? Z_1 + Y_1\} \\ branch_5 &= \{X \approx_? X_1, W \approx_? Z_1 + W_1, \pi X \approx_? W_1 + X_1, \pi Y \approx_? Z_1\} \\ branch_6 &= \{X \approx_? Y_1, W \approx_? Z_1 + W_1, \pi X \approx_? W_1, \pi Y \approx_? Z_1 + Y_1\} \\ branch_7 &= \{X \approx_? Y_1 + X_1, W \approx_? Z_1 + W_1, \pi X \approx_? W_1 + X_1, \pi Y \approx_? Z_1 + Y_1\} \end{aligned}$$

The next step is to instantiate moderated variables. We denote branch i by B_i , the substitution computed in this branch by σ_i and show the result after performing the instantiations. For brevity, when presenting σ_i we omit the instantiation of variables X_1 , W_1 , Y_1 , Z_1 since they were not in the initial problem.

$$B1 - \{\pi X \approx_? X\}, \sigma_1 = \{W \mapsto \pi Y\}$$

$$B2 - \sigma_2 = \{W \mapsto \pi^2 Y, X \mapsto \pi Y\}$$

$$B3 - \{f(\pi^2 Y, \pi X_1) \approx_? f(W, X_1)\}, \sigma_3 = \{X \mapsto f(\pi Y, X_1)\}$$

$$B4 - No \ solution$$

$$B5 - No \ solution$$

$$B6 - \sigma_6 = \{W \mapsto f(Z_1, \pi X), Y \mapsto f(\pi^{-1} Z_1, \pi^{-1} X)\}$$

$$B7 - \{f(\pi Y_1, \pi X_1) \approx_? f(W_1, X_1)\}, \sigma_7 = \{X \mapsto f(Y_1, X_1), W \mapsto f(Z_1, W_1), Y \mapsto f(\pi^{-1} Z_1, \pi^{-1} Y_1)\}$$

Branches 3 and 7 are a renaming of the original problem

$$f(X,W) \approx_{?} f(\pi X, \pi Y).$$

Regarding Branch 3, notice that if we rewrite $\sigma_3 = \{X \mapsto f(\pi Y, X_1)\}$ as $\sigma'_3 = \{Y \mapsto \pi^{-1}Y_1, X \mapsto f(\pi Y, X_1), \}$, then the equational constraint of the mentioned branch is simply:

$$f(X_1, W_1) \approx_? f(\pi X_1, \pi Y_1).$$

Regarding Branch 7, it's even simpler to see the renaming, as the equational constraint is:

$$f(X_1, W_1) \approx_? f(\pi X_1, \pi Y_1),$$

D An Example of ACMATCH

In this section we continue the example of Appendix C and show how ACMATCH would solve the matching equational constraint $f(X,W) \approx_? f(\pi X,\pi Y)$. Algorithm 1 would be called with parameters $\Gamma = \emptyset$, $P = \{f(X,W) \approx_? f(\pi X,\pi Y)\}$, $\sigma = id$, $V = \{X,W,Y\}$ and $\mathcal{X} = \{X,Y\}$. We would enter the **else if** of line 36 and function APPLYACSTEP would be called.

Function APPLYACSTEP would call SOLVEAC, which would find the corresponding linear Diophantine equation and generate a set of solutions to it, associating a new moderated variable to each solution. To simplify our example, we assume that the solutions computed by SOLVEAC and the new moderated variables introduced by it are the ones represented in Table C. In the end, SOLVEAC would construct the new equational constraints and the algorithm would branch. However, since we discard the branches where a protected variable is paired with an AC function application, we do not generate all the 7 branches of the example in Section C, we only generate 2 branches:

$$branch_1 = \{X \approx_? X_1, W \approx_? Z_1, \pi X \approx_? X_1, \pi Y \approx_? Z_1\}$$
$$branch_2 = \{X \approx_? Y_1, W \approx_? W_1, \pi X \approx_? W_1, \pi Y \approx_? Y_1\}$$

Next, APPLYACSTEP will call INSTANTIATESTEP, which would instantiate the moderated variables that it can. When encountering an equation such as $\pi X \approx_? \pi' X$, INSTANTIATESTEP solves it by adding $ds(\pi,\pi')\# X$ to the context Γ we are working with. This is complete because we are doing matching, not unification. We denote each branch i by B_i and, as was done in Section C, when presenting σ_i we omit the instantiation of variables X_1 , W_1 , Y_1 , Z_1 since they were not in the initial problem. The results are shown below:

$$B1 - \Gamma_1 = ds(\pi, id) \# X, \sigma_1 = \{W \mapsto \pi Y\}$$

$$B2 - \Gamma_2 = \emptyset, \sigma_2 = \{W \mapsto \pi^2 Y, X \mapsto \pi Y\}$$

When function APPLYACSTEP finishes, there is no equational constraint left and ACMATCH also returns. There are two solutions computed in this example: $(\Gamma_1, \sigma_1) = (ds(\pi, id) \# X, \{W \mapsto \pi Y\})$ and $(\Gamma_2, \sigma_2) = (\emptyset, \{W \mapsto \pi^2 Y, X \mapsto \pi Y\})$.

E Grammar of Terms and the Need for Well-Formed Terms

Since the formalisation of nominal AC-match extends the formalisation of first-order AC-unification and the latter restricts the terms received by the algorithm to well-formed terms, it is natural that the nominal AC-match formalisation would also enforce this restriction. In this section we review the motives behind this decision

Example 2. Some examples to illustrate the behaviour of $Args_f$.

```
-Args_f(a) = (a). - Args_f(f\langle c, b \rangle) = (c, b). 
-Args_f(\pi \cdot Y) = (\pi \cdot Y). - Args_f(f\langle c, b \rangle) = (a, b, c). - Args_f(g\langle c, b \rangle) = (g\langle c, b \rangle).
```

As mentioned before, terms were defined as shown in Definition 1 in order to make it easier to eventually adapt the formalisation to the nominal setting. However, two issues arose in the formalisation that motivated us to define well-formed terms (Definition 2) and restrict the terms in the unification problem that our algorithm receive to well-formed terms.

The first issue has to do with AC-functions that receive only one argument, something allowed in the grammar of terms. Let f be an AC-function symbol and consider Example 3, which shows that $ff\langle a,b\rangle\approx_? f\langle a,b\rangle$. This is problematic in first-order AC-unification, because it means that $P=\{X\approx_? fX\}$ has a solution, for instance $\sigma=\{X\mapsto f\langle a,b\rangle\}$. However, if the algorithm of [7] received this unification problem P, it would return nil. In the definition of well-formed terms, we avoid this problem by requiring that every AC-function application $f^{AC}s$ that is a subterm of a well-formed term t does not receive only one argument.

Example 3. Let f be an AC-function symbol. Consider the terms $t \equiv ff\langle a, b \rangle$ and $s \equiv f\langle a, b \rangle$. Two AC function applications are equal (modulo AC) if and only if their list of arguments are permutations of each other. In our particular case we have $Args_f(t) = (a, b) = Args_f(s)$ and therefore $t \approx s$.

The second issue is with terms that are pairs. As mentioned before, pairs are to be used inside a term t to encode a tuple of arguments to a function. If t and s are not pairs and $Args_f(t)$ and $Args_f(s)$ are permutations of each other then it is possible to prove that $t \approx s$. This result we just described was used in the proof of completeness of SOLVEAC for first-order AC-unification (see the extended version of [7]), continued being used in the nominal AC-matching formalisation, and is the reason why we imposed that a well-formed term t is not a pair.

Example 4. Let f be an AC-function symbol and g be a syntactic function symbol. The following terms are well-formed terms:

```
 - f\langle a, \langle b, c \rangle \rangle. 
 - f f\langle a, \langle b, c \rangle \rangle \text{ (here } Args_f(f f\langle a, \langle b, c \rangle \rangle) = (a, b, c)). 
 - a. 
 - g(Y).
```

The following terms are not well-formed terms:

 $\begin{array}{l}
- fX. \\
- \langle a, b \rangle.
\end{array}$

E.1 Equal Terms May Not Have the Same Size

A drawback of our grammar of terms is that we can have well-formed terms that are equal modulo AC that do not have the same size. Let f be an AC-function symbol and consider for instance the terms $t \equiv f \langle f(a,b),c \rangle$ and $s \equiv f \langle \langle a,b \rangle,c \rangle$. These terms are equal modulo AC. Indeed $Args_f(t) = (a,b,c) = Args_f(s)$ but according to the definition of size we have size(t) = 7 and size(s) = 6. An alternative definition of size, called $size_2$, that has this property (Theorem 15) is given below.

Definition 15 ($size_2$ \bigcirc). We define the $size_2$ of a term t recursively as follows:

```
\begin{array}{l} - \ size_2(a) = 1 \\ - \ size_2(Y) = 1 \\ - \ size_2(\langle \rangle) = 1 \\ - \ size_2(\langle t_1, t_2 \rangle) = size_2(t_1) + size_2(t_2) \\ - \ size_2(ft_1) = 1 + size_2(t_1) \\ - \ size_2(f^{AC}t_1) = \sum_{t_i \in Args_f(f^{AC}t_1)} size_2(t_i) \end{array}
```